

Eddie Aguilar

## Quiz - Forecasting and beyond

Try to answer each question or exercise by yourself and without any external help. If needed, you can check your class notes, exercises, documentation online, or even your peers.

Let's see what you've learned regarding the forecasting workflow so far. Answer the following questions to the best of your abilities.

**What are the type of components that a time series can exhibit?**

- ☒ A residual component.
- ☒ A trend-cycle component.
- ☐ A fluctuating component.
- ☐ Autocorrelation component.
- ☒ Seasonal component(s).

Correct!

**What is the purpose of performing calendar, population, or inflation adjustments to a time series?**

- ☐ To remove the trend from it.
- ☒ Reduce/remove variation that is not relevant to the forecast/analysis done.
- ☐ Correct possible heteroscedasticity in the ts.
- ☐ Make the time series look prettier.

Correct!

R Code [Start Over](#)

[Run Code](#)

```
1 # This is here just in case you need to run any R code.  
2  
3
```

[Next Topic](#)

## Forecasting - Prerequisites

Let's do some exercises and practice what we've seen so far. We will download some data and go through all the forecasting workflow. To start, load the necessary packages into the environment.

```
R Code Start Over Hint Run Code
1 library(tidyverse)
2 library(fpp3)
3
```

### ✓ Importing data

To simplify the data importing step, we will import a time series directly from the [FRED](#) website using the `td_get()` function from the **tidyquant** package. In this case we will work with the Advance Retail Sales, specifically sales from Food and Beverages stores (`RSDBSN`). Go ahead and load it and save it to a variable called `ex1_tbl`. Some code is already provided for you.

```
R Code Start Over Hints Run Code
1 ex1_tbl <- tidyquant::td_get(
2   x = "RSDBSN",
3   get = "economic.data",
4   from = "1992-01-01"
5 )
```

### ✓ Wrangling the data

Inspect the data frame you just downloaded using function `class()`.

```
R Code Start Over Hint Run Code
1 class(ex1_tbl)
2
3
```

```
[1] "tbl_df"      "tbl"        "data.frame"
```

You should see something like:

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

The first argument `tbl_df` indicates we have a **tibble**. We must convert it to a **tsibble**. But first, we need to make sure our **date** variable is in the correct format for your time series. Change it if needed and overwrite `ex1_tbl`, otherwise leave it as is.

```
R Code Start Over Hint Run Code
1
2 ex1_tbl <- ex1_tbl |>
3   as_tsibble(index = date)
4   ex1_tbl
5
```

| symbol<br><chr> | date<br><date> | price<br><int> |
|-----------------|----------------|----------------|
| RSDBSN          | 1995-05-01     | 33178          |
| RSDBSN          | 1995-06-01     | 33060          |
| RSDBSN          | 1995-07-01     | 33520          |

Now you can convert your `tibble` into a `tsibble`. Do so using `as_tsibble()` and store your result now in a variable named `ex1_tsbl`.

Any given `tsibble` needs an *index*, which is the date variable. When do you need to specify the `key` argument for the `tsibble` to work?

✓ When you data frame is in a *long* format, and you have more than one time series in the data frame.

X Always.

Correct!

```
R Code Start Over Hint Run Code
1 ex1_tsbl <- ex1_tbl |>
2   as_tsibble(index = date)
3 ex1_tsbl
4
5
```

| symbol<br><chr> | date<br><mth> | price<br><int> |
|-----------------|---------------|----------------|
| RSDBSN          | 1992 Jan      | 29589          |
| RSDBSN          | 1992 Feb      | 28570          |
| RSDBSN          | 1992 Mar      | 29682          |
| RSDBSN          | 1992 Apr      | 30228          |
| RSDBSN          | 1992 May      | 31677          |
| RSDBSN          | 1992 Jun      | 30769          |
| RSDBSN          | 1992 Jul      | 32402          |
| RSDBSN          | 1992 Aug      | 31469          |
| RSDBSN          | 1992 Sep      | 30162          |
| RSDBSN          | 1992 Oct      | 31407          |

1-10 of 375 rows

Previous [1](#) [2](#) [3](#) [4](#) [5](#) [6](#) ... [38](#) [Next](#)

Check now the class of this new variable `ex1_tsbl`:

```
R Code Start Over Solution Run Code
1 class(ex1_tsbl)
2
3
```

```
[1] "tbl_ts"      "tbl_df"      "tbl"        "data.frame"
```

It should now display as its first argument `tbl_ts`, which indicates the object is a `tsibble`.

## ✓ Train/test splits

Split your data set so that you have the last **2 years** as test data. Save your training set to a variable called `ex1_train`.

R Code

Start Over

Hint

Run Code

```
1 ex1_train <- filter_index(ex1_tsbl, "2005 Jan" ~ "2021 Dec")
2 ex1_train
```

| symbol | date     | price |
|--------|----------|-------|
| <chr>  | <mth>    | <int> |
| RSDBSN | 2021 Sep | 73851 |
| RSDBSN | 2021 Oct | 76702 |
| RSDBSN | 2021 Nov | 76532 |
| RSDBSN | 2021 Dec | 83523 |
|        |          |       |
|        |          |       |
|        |          |       |
|        |          |       |
|        |          |       |
|        |          |       |

201-204 of 204 rows

Previous

1

...

16

17

18

19

20

21

Next

[Previous Topic](#)[Next Topic](#)

# Exploratory Data Analysis


## ✓ Visualization

Time plot

Plot your training set using `autoplot()`.

```
R Code Start Over Run Code  
1 ex1_train <- filter_index(ex1_tsbl, "2005 Jan" ~ "2021 Dec")  
2 autoplot(ex1_train)  
3  
4
```

Plot variable not specified, automatically selected ``.vars = price``



Which patterns do you see on it?

Seasonality

✓ Yes

X No

Correct!

Trend

✓ Yes

X No

Correct!

Constant variance (Homoscedasticity)

X Yes

✓ No

Correct!

## Mathematical Transformations


It seems that we may need a transformation to stabilize the variance of the ts. Try using a log transformation and plot it to verify if it works:

R Code

[Start Over](#)

[Run Code](#)

```
1 autoplot(ex1_train, log(price))  
2  
3
```




If you decided a transformation was in order, remember to use it from here on over.

## Season plot

Now let's check a **season plot**, to get a closer look about the seasonality.

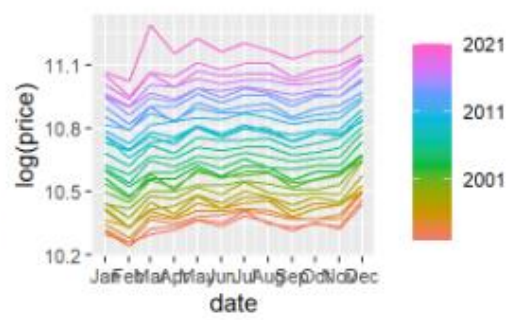


```
R Code Start Over Hint Run Code  
1 gg_season(ex1_train, log(price))  
2  
3
```



There are many other types of plots that could aid you in getting more insights from the data. Feel free to try them here:

```
R Code Start Over Run Code  
1  
2 gg_tsdisplay(ex1_train, log(price))  
3
```





## Modeling the time series

So far, we only know the basic benchmark models (*Mean, Naïve, Seasonal Naïve, and Drift* methods). Based upon what you've discovered from this time series so far, create a `mable` (table of models) estimating those which you consider will fit better to the data. Store the `mable` in `ex1_fit`.

If you chose to use a transformation, remember to apply it inside the model specification, so that R can detect it.

R Code

Start Over

Hints

Run Code

```
1 ex1_fit <- ex1_train |>
2   model()
3   snaive = SNAIVE(log(price)),
4   drift = RW(log(price) ~ drift())
5
6 ex1_fit
```

| symbol | snaive    | drift     |
|--------|-----------|-----------|
| <chr>  | <lst_md1> | <lst_md1> |
| RSDBSN | <lst_md1> | <lst_md1> |

1 row

### ✓ Model evaluation

How can you tell if your model(s) fit the data correctly? We can do so in several ways:

- Checking the residual diagnostics
  - Visually
  - Using statistical tests, such as the Ljung-Box test for autocorrelation.
- Verify the error metrics on the training set.

### Quiz

We expect the residuals of the model to be entirely white noise. A variable can be considered white noise if:

- ☐ Its variance is 1.
- ☒ It has no seasonality
- ☒ It has no significant autocorrelation.
- ☒ It exhibits no trend.

Correct!

When referring to the residuals of a model, besides what it was said on the previous question, we also check that:

- ✓ The mean is close to zero.
- ✓ They are normally distributed.
- ✓ They present homoscedasticity (time-invariant variance).
- ✗ The variance is close to one.

Correct!

## ✓ Residual diagnostics


Do the residual diagnostics with the help of the `gg_tseriesiduals()` function.


```
R Code Start Over Hint Run Code  
1 ex1_fit <- ex1_train |>  
2   model(  
3     snaiive = SNAIVE(log(price)),  
4     drift = RW(log(price) ~ drift())  
5   )  
6  
7 gg_tseriesiduals(select(ex1_fit, snaiive))  
8 gg_tseriesiduals(select(ex1_fit, drift))  
9
```

Warning: Removed 12 rows containing missing values (`geom\_line()`).

Warning: Removed 12 rows containing missing values (`geom\_point()`).

Warning: Removed 12 rows containing non-finite values (`stat\_bin()`).





## ✓ Portmanteau tests

To further assess whether the model(s) provided are producing white noise residuals, test it out using Ljung-Box.

| R Code <a href="#">Start Over</a> <a href="#">Hints</a> <a href="#">Run Code</a>                  |                 |                  |                    |
|---|-----------------|------------------|--------------------|
| <pre>1 ex1_fit  &gt; 2   augment()  &gt; 3   features(.resid, ljung_box, lag = 24, dof = 0)</pre> |                 |                  |                    |
| symbol<br><chr>   | .model<br><chr> | lb_stat<br><dbl> | lb_pvalue<br><dbl> |
| RSDBSN  | drift           | 885.0620         | 0                  |
| RSDBSN  | seasonal_naive  | 630.6089         | 0                  |

2 rows

## ✓ Training Accuracy

Use the `accuracy()` function together with the `mable` to see the different error metrics calculated on the training set.

| R Code <a href="#">Start Over</a> <a href="#">Hints</a> <a href="#">Run Code</a> |                 |                |             |               |
|--|-----------------|----------------|-------------|---------------|
| <pre>1 accuracy(ex1_fit) 2 3</pre>   |                 |                |             |               |
| symbol<br><chr>  | .model<br><chr> | .type<br><chr> | ME<br><dbl> | RMSE<br><dbl> |
| RSDBSN   | seasonal_naive  | Training       | 1441.94362  | 2082.457      |
| RSDBSN   | drift           | Training       | 3.69474     | 3053.115      |

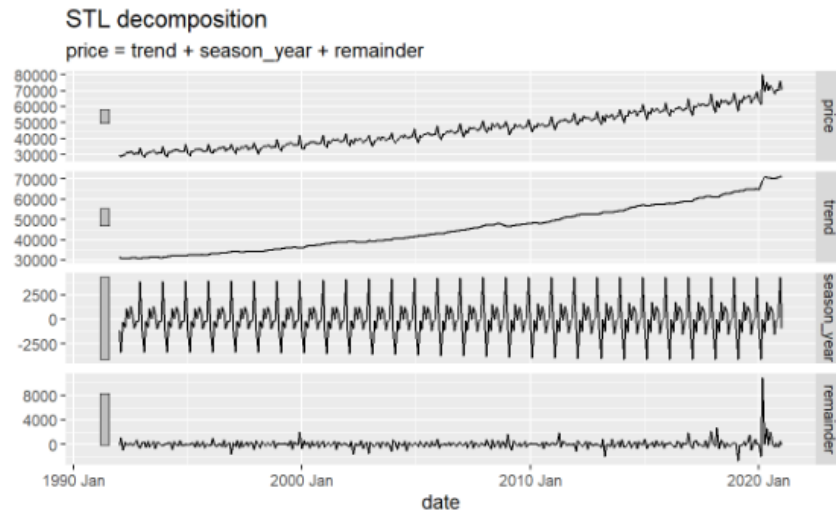
2 rows | 1-5 of 11 columns

## ✓ Model with decomposition

Let's add another model to our analysis, using a STL decomposition and model each component independently. Save this new model in

`ex1_fit_dcmp`:

```
R Code Start Over Hints Run Code
1 ex1_fit_dcmp <- ex1_train %>%
2   model(STL(price ~ trend(window = 7) + season(window=21), robust=TRUE)) %>%
3   components()
4
5 autoplot(ex1_fit_dcmp)
6
7
```



Now join this new model with our original `mable` using `left_join()`:

```
R Code Start Over Run Code
1 (ex1_fit <- left_join(x = ex1_fit, y = ex1_fit_dcmp))
2
3
4
```

Joining with `by = join\_by(symbol)`

| symbol | seasonal_naive | drift  | stlf   |
|--------|----------------|--------|--------|
| <chr>  | <lstm>         | <lstm> | <lstm> |
| RSDBSN | <lstm>         | <lstm> | <lstm> |
| 1 row  |                |        |        |

Previous Topic

Next Topic

# Forecasting on the test set

## ✓ Getting the `fable`

Produce forecasts for the test set. Store the `fable` (forecast table) in `ex1_fc`.

R Code

Start Over

Hints

Run Code

```
1 (ex1_fc <- forecast(ex1_fit, h = 16))
2
3
```

| symbol | .model | date     | price  | .mean    |
|--------|--------|----------|--------|----------|
| <chr>  | <chr>  | <month>  | <dist> | <dbl>    |
| RSDBSN | stlf   | 2021 Oct | <dist> | 73260.16 |
| RSDBSN | stlf   | 2021 Nov | <dist> | 73997.00 |
| RSDBSN | stlf   | 2021 Dec | <dist> | 79738.74 |
| RSDBSN | stlf   | 2022 Jan | <dist> | 72788.18 |
| RSDBSN | stlf   | 2022 Feb | <dist> | 67893.19 |
| RSDBSN | stlf   | 2022 Mar | <dist> | 74677.21 |
| RSDBSN | stlf   | 2022 Apr | <dist> | 72818.65 |
| RSDBSN | stlf   | 2022 May | <dist> | 77448.55 |

41-48 of 48 rows

Previous12345Next

## ✓ Forecast plots

Plot the forecasts produced by each model, along with the training and test set. Tip: you'll need the `fable`, and the full `tsibble`.

R Code

Start Over

Solution

Run Code

```
1 ex1_fc |>
2 autoplot(filter_index(ex1_tsbl, "2005" ~ .), level = FALSE) + |
3 ggtitle("Forecast")
4
5
```

Warning: There was 1 warning in `filter()`.  
i In argument: `time\_in(date, ...)`.  
Caused by warning:  
! `yearmonth()` may yield unexpected results.  
i Please use arg `format` to supply formats.

Warning: Plot argument `level` should be a numeric vector of levels to display.  
Setting `level = NULL` will remove the intervals from the plot.

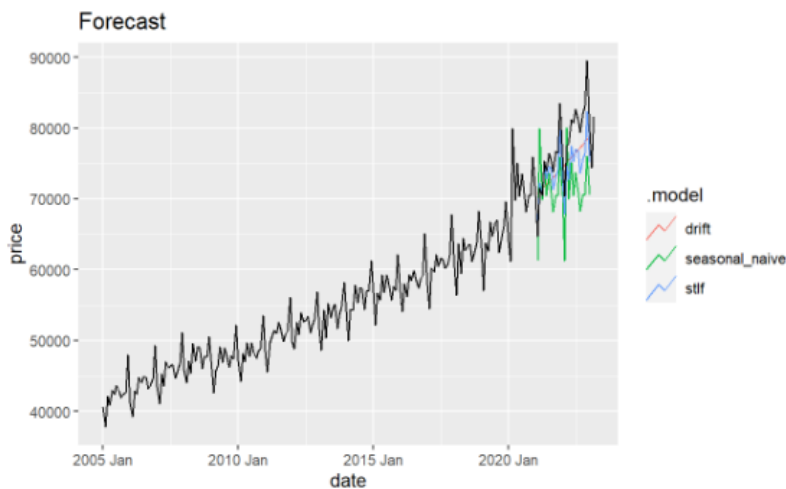
## ✓ Forecast plots

Plot the forecasts produced by each model, along with the training and test set. Tip: you'll need the `fable`, and the full `tsibble`.

```
R Code Start Over Solution Run Code
1 ex1_fc |>
2 autoplot(filter_index(ex1_tsbl, "2005" ~ .), level = FALSE) +
3 ggtitle("Forecast")
4
5
```

Warning: There was 1 warning in `filter()``.  
i In argument: `time\_in(date, ...)`.  
Caused by warning:  
! `yearmonth()` may yield unexpected results.  
i Please use arg `format` to supply formats.

Warning: Plot argument `level` should be a numeric vector of levels to display.  
Setting `level = NULL` will remove the intervals from the plot.

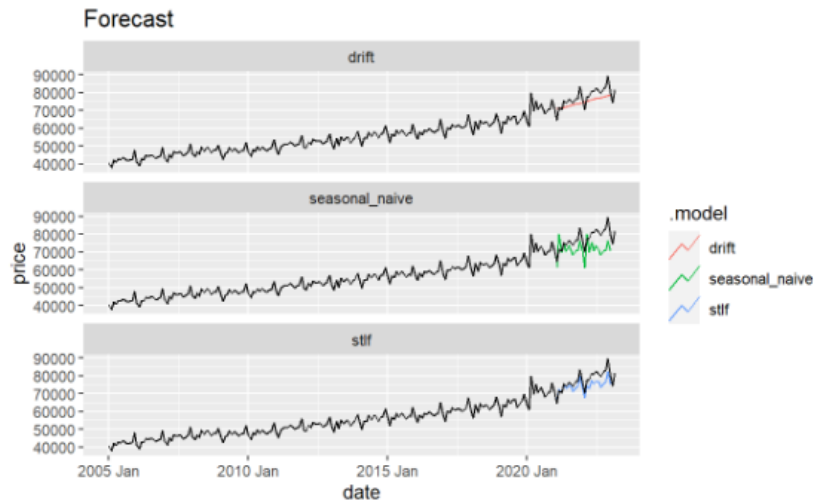


Try splitting each forecast in its own subplot, defining free scales for the *y* axis:

```
R Code Start Over Hints Run Code
1 ex1_fc |>
2 autoplot(filter_index(ex1_tsbl, "2005" ~ .), level = FALSE) +
3 ggtitle("Forecast") +
4 facet_wrap(~ .model, scales = "free_y", ncol = 1)
```

Warning: There was 1 warning in `filter()`.  
 i In argument: `time\_in(date, ...)`.  
 Caused by warning:  
 ! `yearmonth()` may yield unexpected results.  
 i Please use arg `format` to supply formats.

Warning: Plot argument `level` should be a numeric vector of levels to display.  
 Setting `level = NULL` will remove the intervals from the plot.



## ✓ Forecasting errors

Estimate the forecast errors on the test set using the `accuracy()` function. You now need to provide the `table`, and the full `tsibble`.

R Code [Start Over](#) [Hints](#) [Run Code](#)

```
1 ex1_fc |>
2   accuracy()ex1_tsbl[]
3
```

| .model<br><chr> | symbol<br><chr> | .type<br><chr> | ME<br><dbl> | RMSE<br><dbl> | MAE<br><dbl> |
|-----------------|-----------------|----------------|-------------|---------------|--------------|
| drift           | RSDBSN          | Test           | 2796.643    | 4543.413      | 3751.480     |
| seasonal_naive  | RSDBSN          | Test           | 6034.057    | 7817.156      | 6953.326     |
| stlf            | RSDBSN          | Test           | 3249.102    | 3967.249      | 3394.971     |

3 rows | 1-6 of 11 columns

[Previous Topic](#)

[Next Topic](#)



## Forecasting the future

### ✓ Refitting the best model

It seems that the model that produces the best forecasts is the one with decomposition (`stlf`). Refit this model using the whole data (not just the training set). Store your newly fit model into `ex1_fit_fut`

```
R Code Start Over Hint Run Code
1 |
2 |
3 | ex1_fit_fut <- model(ex1_tsbl, STL(price ~ trend(window = 7) + season(window=21), robust=TRUE)) %>%
4 |   components()
5 |
6 |
```

### ✓ Forecast the future

Produce a 2-year forecast into the future. Save it to `ex1_fc_fut`

```
R Code Start Over Run Code
1 | ex1_fc_fut <- ex1_fit_fut |>
2 |   forecast(h = "2 years")
3 |
```

### ✓ Forecast plot

Plot the final forecast.

```
R Code Start Over Solution Run Code
1 | autoplot(ex1_fc_fut, ex1_tsbl)
2 |
3 |
```

